

SUBMISSION DATE: 21.09.2023

ASSIGNMENT – 4

Artificial Intelligence & Machine Learning in collaboration
with Google (Applied Data Science)

Name: Diya Maria Sanjay

Regno: 21BML0113

Branch: Btech ECE with specialization in Biomedical
Engineering

Campus: VIT Vellore

#Grapes to Greatness: Machine Learning in Wine Quality Prediction

0.0.1 Task 1 : Load the Dataset

```
[1]: # import required libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df = pd.read_csv('/content/winequality-red.csv')
df.head()
```

```
[2]: fixed_acidity volatile_acidity citric_acid residual_sugar chlorides \
0          7.4          0.70          0.00          1.9          0.076
1          7.8          0.88          0.00          2.6          0.098
2          7.8          0.76          0.04          2.3          0.092
3         11.2          0.28          0.56          1.9          0.075
4          7.4          0.70          0.00          1.9          0.076
free_sulfur_dioxide total_sulfur_dioxide density pH sulphates \
0          11.0    34.0 0.9978 3.51 0.56
1          25.0    67.0 0.9968 3.20 0.68
2          15.0    54.0 0.9970 3.26 0.65
3          17.0    60.0 0.9980 3.16 0.58
4          11.0    34.0 0.9978 3.51 0.56
```

```
alcohol quality
0    9.4    5 1
    9.8    5 2
    9.8    5 3
    9.8    6
4    9.4    5
```

0.0.2 Task 2 : Data preprocessing including visualization

```
[3]: df.shape
```

```
[3]: (1599, 12)
```

```
[4]: df.info()
```

```
<class
'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to
1598 Data columns (total 12
columns):
#      Column              Non-Null Count  Dtype
---  -
-----
```

```

0  fixed_acidity      1599 non-null float64
1  volatile_acidity   1599 non-null float64
2  citric_acid        1599 non-null float64
3  residual_sugar     1599 non-null float64
4  chlorides          1599 non-null float64
5  free_sulfur_dioxide 1599 non-null float64
6  total_sulfur_dioxide 1599 non-null float64
7  density            1599 non-null float64
8  pH                1599 non-null float64
9  sulphates          1599 non-null float64
10 alcohol            1599 non-null float64
11 quality            1599 non-null int64

```

```
dtypes: float64(11), int64(1)
```

```
memory usage: 150.0 KB
```

```
[5]: df.isnull().sum() # There are no null values in the dataset.
```

```

[5]: fixed_acidity      0
     volatile_acidity   0
     citric_acid        0
     residual_sugar     0
     chlorides          0
     free_sulfur_dioxide 0
     total_sulfur_dioxide 0
     density            0
     pH                0
     sulphates          0
     alcohol            0
     quality            0
     dtype:
     int64

```

```
[6]: df.describe() # Descriptive Statistics
```

```

[6]:      fixed_acidity      volatile_acidity      citric_acid
      residual_sugar \
count    1599.000000    1599.000000  1599.000000    1599.000000
mean         8.319637         0.527821    0.270976     2.538806
std          1.741096         0.179060    0.194801     1.409928
min           4.600000         0.120000    0.000000     0.900000
25%           7.100000         0.390000    0.090000     1.900000
50%           7.900000         0.520000    0.260000     2.200000
75%           9.200000         0.640000    0.420000     2.600000
max          15.900000         1.580000    1.000000    15.500000

      chlorides free_sulfur_dioxide total_sulfur_dioxide  density \
count    1599.000000    1599.000000    1599.000000  1599.000000
mean         0.087467        15.874922        46.467792    0.996747
std          0.047065        10.460157        32.895324    0.001887
min           0.012000         1.000000         6.000000    0.990070
25%           0.070000         7.000000        22.000000    0.995600

```

50%	0.079000	14.000000	38.000000	0.996750
75%	0.090000	21.000000	62.000000	0.997835
max	0.611000	72.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	
1599.000000	mean	3.311113	0.658149	10.422983
5.636023	std	0.154386	0.169507	1.065668
min	2.740000	0.330000	8.400000	3.000000
3.210000	0.550000	9.500000	5.000000	
50%	3.310000	0.620000	10.200000	6.000000
	3.400000	0.730000	11.100000	6.000000
	4.010000	2.000000	14.900000	8.000000

```
[7]: df.corr()
```

```
[7]:
```

	fixed_acidity	volatile_acidity		
	citric_acid \			
fixed_acidity	1.000000	-0.256131	0.671703	
volatile_acidity	-0.256131	1.000000	-0.552496	
citric_acid	0.671703	-0.552496	1.000000	
residual_sugar	0.114777	0.001918	0.143577	
chlorides	0.093705	0.061298	0.203823	
free_sulfur_dioxide	-0.153794	-0.010504	-0.060978	
total_sulfur_dioxide	-0.113181	0.076470	0.035533	
density	0.668047	0.022026	0.364947	
pH	-0.682978	0.234937	-0.541904	
sulphates	0.183006	-0.260987	0.312770	
alcohol	-0.061668	-0.202288	0.109903	
quality	0.124052	-0.390558	0.226373	
	residual_sugar	chlorides	free_sulfur_dioxide \	
fixed_acidity	0.114777	0.093705	-0.153794	
volatile_acidity	0.001918	0.061298	-0.010504	
citric_acid	0.143577	0.203823	-0.060978	
residual_sugar	1.000000	0.055610	0.187049	
chlorides	0.055610	1.000000	0.005562	
free_sulfur_dioxide	0.187049	0.005562	1.000000	
total_sulfur_dioxide	0.203028	0.047400	0.667666	
density	0.355283	0.200632	-0.021946	
pH	-0.085652	-0.265026	0.070377	
sulphates	0.005527	0.371260	0.051658	
alcohol	0.042075	-0.221141	-0.069408	
quality	0.013732	-0.128907	-0.050656	
	total_sulfur_dioxide	density	pH	sulphates
fixed_acidity	-0.113181	0.668047	-0.682978	0.183006

volatile_acidity	0.076470	0.022026	0.234937	-
	0.260987			
citric_acid	0.035533	0.364947	-0.541904	0.312770
residual_sugar	0.203028	0.355283	-0.085652	0.005527
chlorides	0.047400	0.200632	-0.265026	0.371260
free_sulfur_dioxide	0.667666	-0.021946	0.070377	0.051658
total_sulfur_dioxide	1.000000	0.071269	-0.066495	0.042947
density	0.071269	1.000000	-0.341699	0.148506
pH	-0.066495	-0.341699	1.000000	-
	0.196648			
sulphates	0.042947	0.148506	-0.196648	1.000000
alcohol	-0.205654	-0.496180	0.205633	0.093595
quality	-0.185100	-0.174919	-0.057731	0.251397

	alcohol	quality
fixed_acidity	-0.061668	0.124052
volatile_acidity	-0.202288	-0.390558
citric_acid	0.109903	0.226373
residual_sugar	0.042075	0.013732
chlorides	-0.221141	-0.128907
free_sulfur_dioxide	-0.069408	-0.050656
total_sulfur_dioxide	-0.205654	-0.185100
density	-0.496180	-
	0.174919	
pH	0.205633	-
	0.057731	
sulphates	0.093595	
	0.251397	
alcohol	1.000000	
	0.476166	
quality	0.476166	
	1.000000	

```
[8]: # Correlation of dependent varriables with the target variable
df.corr().quality.sort_values(ascending = False)
```

```
[8]: quality          1.000000
     alcohol         0.476166
     sulphates       0.251397
     citric_acid      0.226373
     fixed_acidity     0.124052
     residual_sugar    0.013732
     free_sulfur_dioxide -
                        0.050656
     pH              -
                        0.057731
     chlorides        -
                        0.128907
```

```
density          -  
                  0.174919  
total_sulfur_dioxide-  
                  0.185100  
volatile_acidity -  
                  0.390558  
Name: quality, dtype: float64
```

Univariate Analysis

```
[9]: sns.distplot(df.sulphates)
```

```
<ipython-input-9-8b271c44c149>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn  
v0.14.0.
```


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df.sulphates)
```

```
[9]: <Axes: xlabel='sulphates', ylabel='Density'>
```



```
[10]: sns.displot(df.chlorides)
```


```
[10]: <seaborn.axisgrid.FacetGrid at 0x7ddd8a543160>
```



Bivariate Analysis


```
[11]: sns.lineplot(x=df.sulphates, y=df.pH)
```

```
[11]: <Axes: xlabel='sulphates', ylabel='pH'>
```

```
[12]: sns.scatterplot(x=df.alcohol, y=df.pH)
```

```
[12]: <Axes: xlabel='alcohol', ylabel='pH'>
```

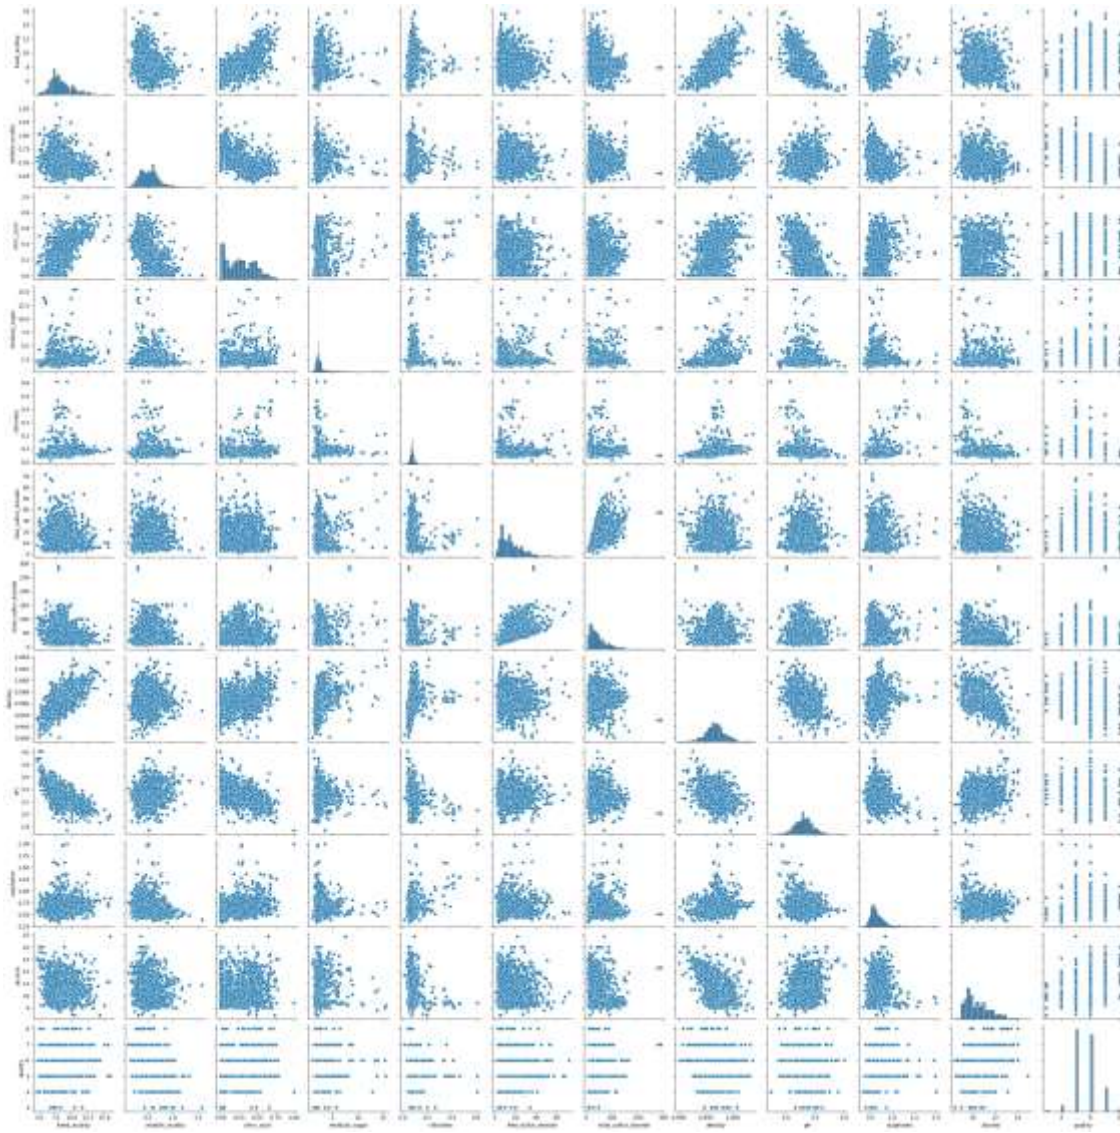


Multivariate Analysis

[13]:


```
sns.pairplot(df)
```

[13]: <seaborn.axisgrid.PairGrid at 0x7ddd4f583280>



```
[14]: # Correlation Heatmap
sns.heatmap(df.corr(),annot=True)
```

```
[14]: <Axes: >
```



Outlier Detection and removal by percentile method & IQR METHOD

```
[16]: df.head()
```

```
[16]: fixed_acidity volatile_acidity citric_acid residual_sugar chlorides \
0          7.4          0.70          0.00          1.9          0.076
1          7.8          0.88          0.00          2.6          0.098
2          7.8          0.76          0.04          2.3          0.092
3         11.2          0.28          0.56          1.9          0.075
4          7.4          0.70          0.00          1.9          0.076
free_sulfur_dioxide total_sulfur_dioxide density pH sulphates \
0          11.0          34.0 0.9978 3.51          0.56
1          25.0          67.0 0.9968 3.20          0.68
2          15.0          54.0 0.9970 3.26          0.65
3          17.0          60.0 0.9980 3.16          0.58
4          11.0          34.0 0.9978 3.51          0.56
alcohol quality
0          9.4          5
```

1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5


```
[49]: # Removing outliers from fixed_acidity column
```

```
f1 = df.fixed_acidity.quantile(0.25) #Q1
f3 = df.fixed_acidity.quantile(0.75) #Q3
IQR_f = f3 - f1
upper_limit_f = f3+(1.5)*(IQR_f)
lower_limit_f = f1-(1.5)*(IQR_f)
print(f1)
print(f3)
print(IQR_f)
print(upper_limit_f)
print(lower_limit_f)
```

```
7.1
8.9
1.8000000000000007
11.600000000000001
4.399999999999999
```

```
[51]: df=df[(df.fixed_acidity<upper_limit_f) &
(df.fixed_acidity>lower_limit_f)] sns.boxplot(df.fixed_acidity)
```

```
[51]: <Axes: >
```




```
[47]: fa_01=df.fixed_acidity.quantile(0.01)
      fa_9=df.fixed_acidity.quantile(0.98)
      print(fa_01)
      print(fa_98)
```

```
5.6
11.6
```

```
[48]: df=df[(df.fixed_acidity>=fa_01) & (df.fixed_acidity<=fa_98)]
      sns.boxplot(df.fixed_acidity)
```

```
[48]: <Axes: >
```



```
[22]: # Removing outliers from volatile_acidity column
```

```
v1 = df.volatile_acidity.quantile(0.25) #Q1
v3 = df.volatile_acidity.quantile(0.75) #Q3
IQR_v = v3 - v1
upper_limit_v = v3+(1.5)*(IQR_v)
lower_limit_v = v1-(1.5)*(IQR_v)
print(v1)
print(v3)
print(IQR_v)
print(upper_limit_v)
print(lower_limit_v)
```

```
0.3925
```

```
0.64
```


```
0.2475
```

```
1.01125
```

```
0.0212500000000000047
```

```
[23]: df=df[(df.volatile_acidity<upper_limit_v) &
(df.volatile_acidity>lower_limit_v)] sns.boxplot(df.volatile_acidity)
```

[23]: <Axes: >



[24]: *# Removing outliers from citric_acid column*

```
c1 = df.citric_acid.quantile(0.25) #Q1
c3 = df.citric_acid.quantile(0.75) #Q3
IQR_c = c3 - c1
upper_limit_c = c3+(1.5)*(IQR_c)
lower_limit_c = c1-(1.5)*(IQR_c)
print(c1)
print(c3)
print(IQR_c)
print(upper_limit_c)
print(lower_limit_c)
```

0.09

0.41


0.31999999999999995

0.8899999999999999

-0.3899999999999999

[25]: `df=df[(df.citric_acid<upper_limit_c) & (df.citric_acid>lower_limit_c)] sns.boxplot(df.citric_acid)`

[25]: <Axes: >




[26]: *# Removing outliers from residual_sugar column*

```
r1 = df.residual_sugar.quantile(0.25) #Q1
r3 = df.residual_sugar.quantile(0.75) #Q3
IQR_r = r3 - r1
upper_limit_r = r3+(1.5)*(IQR_r)
lower_limit_r = r1-(1.5)*(IQR_r)
print(r1)
print(r3)
print(IQR_r)
print(upper_limit_r)
print(lower_limit_r)
```

```
1.9
2.6
0.7000000000000002
3.6500000000000004
0.8499999999999996
```

[27]: `df=df[(df.residual_sugar<upper_limit_r) & (df.residual_sugar>lower_limit_r)]` `sns.boxplot(df.residual_sugar)`

[27]: <Axes: >




```
[34]: rs_02=df.residual_sugar.quantile(0.02)
      rs_96=df.residual_sugar.quantile(0.96)
      print(rs_02)
      print(rs_96)
```

```
1.4
3.01599999999999854
```

```
[35]: df=df[(df.residual_sugar>=rs_02) & (df.residual_sugar<=rs_96)]
      sns.boxplot(df.residual_sugar)
```

[35]: <Axes: >




```
[36]: # Removing outliers from chlorides column
```

```
ch1 = df.chlorides.quantile(0.25) #Q1
ch3 = df.chlorides.quantile(0.75) #Q3
IQR_ch = ch3 - ch1
upper_limit_ch = ch3+(1.5)*(IQR_ch)
lower_limit_ch = ch1-(1.5)*(IQR_ch)
print(ch1)
print(ch3)
print(IQR_ch)
print(upper_limit_ch)
print(lower_limit_ch)
```

```
0.07
0.089
0.018999999999999999
0.11749999999999998
0.041500000000000002
```

```
[37]: df=df[(df.chlorides<upper_limit_ch) & (df.chlorides>lower_limit_ch)]
sns.boxplot(df.chlorides)
```

```
[37]: <Axes: >
```




```
[44]: ch_01=df.chlorides.quantile(0.01)
      ch_97=df.chlorides.quantile(0.97)
      print(ch_01)
      print(ch_97)
```

```
0.049890000000000004
0.106
```

```
[45]: df=df[(df.chlorides>=ch_01) & (df.chlorides<=ch_97)]
      sns.boxplot(df.chlorides)
```

```
[45]: <Axes: >
```




```
[52]: # Removing outliers from free_sulfur_dioxide column
```

```
fs1 = df.free_sulfur_dioxide.quantile(0.25) #Q1
fs3 = df.free_sulfur_dioxide.quantile(0.75) #Q3
IQR_fs = fs3 - fs1
upper_limit_fs = fs3+(1.5)*(IQR_fs)
lower_limit_fs = fs1-(1.5)*(IQR_fs)
print(fs1)
print(fs3)
print(IQR_fs)
print(upper_limit_fs)
print(lower_limit_fs)
```

```
8.0
21.0
13.0
40.5 -
11.5
```

```
[53]: df=df[(df.free_sulfur_dioxide<upper_limit_fs) & (df.
    ↪free_sulfur_dioxide>lower_limit_fs)]
sns.boxplot(df.free_sulfur_dioxide) [53]: <Axes: >
```




```
[54]: # Removing outliers from total_sulfur_dioxide column
```

```
ts1 = df.total_sulfur_dioxide.quantile(0.25) #Q1
ts3 = df.total_sulfur_dioxide.quantile(0.75) #Q3
IQR_ts = ts3 - ts1
upper_limit_ts = ts3+(1.5)*(IQR_ts)
lower_limit_ts = ts1-(1.5)*(IQR_ts)
print(ts1)
print(ts3)
print(IQR_ts)
print(upper_limit_ts)
print(lower_limit_ts)
```

```
23.0
57.0
34.0
108.0
-28.0
```

```
[55]: df=df[(df.total_sulfur_dioxide<upper_limit_ts) & (df.
    total_sulfur_dioxide>lower_limit_ts)]
sns.boxplot(df.total_sulfur_dioxide)
```

[55]: <Axes: >




```
[60]: ts_01=df.total_sulfur_dioxide.quantile(0.01)
      ts_97=df.total_sulfur_dioxide.quantile(0.97)
      print(ts_01)
      print(ts_97)
```

```
9.0
89.0
```

```
[61]: df=df[(df.total_sulfur_dioxide>=ts_01) &
      (df.total_sulfur_dioxide<=ts_97)] sns.boxplot(df.total_sulfur_dioxide)
```

[61]: <Axes: >




```
[62]: # Removing outliers from density column
```

```
d1 = df.density.quantile(0.25) #Q1
d3 = df.density.quantile(0.75) #Q3
IQR_d = d3 - d1
upper_limit_d = d3+(1.5)*(IQR_d)
lower_limit_d = d1-(1.5)*(IQR_d)
print(d1)
print(d3)
print(IQR_d)
print(upper_limit_d)
print(lower_limit_d)
```

```
0.9955
0.99745
0.0019499999999998963
1.0003749999999998
0.9925750000000002
```

```
[63]: df=df[(df.density<upper_limit_d) & (df.density>lower_limit_d)]
sns.boxplot(df.density)
```

```
[63]: <Axes: >
```





```
[64]: d_01=df.density.quantile(0.01)
      d_99=df.density.quantile(0.99)
      print(d_01)
      print(d_99)
```

```
0.9933132
1.0
```

```
[65]: df=df[(df.density>=d_01) & (df.density<=d_99)]
      sns.boxplot(df.density)
```

```
[65]: <Axes: >
```




```
[66]: # Removing outliers from pH column

pH1 = df.pH.quantile(0.25) #Q1
pH3 = df.pH.quantile(0.75) #Q3
IQR_pH = pH3 - pH1
upper_limit_pH = pH3+(1.5)*(IQR_pH)
lower_limit_pH = pH1-(1.5)*(IQR_pH)
print(pH1)
print(pH3)
print(IQR_pH)
print(upper_limit_pH)
print(lower_limit_pH)
```

```
3.2425
3.41
0.16749999999999998
3.66125
2.99125
```

```
[67]: df=df[(df.pH<upper_limit_pH) & (df.pH>lower_limit_pH)]
sns.boxplot(df.pH)
```

```
[67]: <Axes: >
```




```
[68]: pH_01=df.pH.quantile(0.01)
      pH_99=df.pH.quantile(0.99)
      print(pH_01)
      print(pH_99)
```

```
3.06
3.6066
```

```
[69]: df=df[(df.pH>=pH_01) & (df.pH<=pH_99)]
      sns.boxplot(df.pH)
```

```
[69]: <Axes: >
```




```
[74]: # Removing outliers from fixed_acidity column
```

```
su_01=df.sulphates.quantile(0.01)
su_98=df.sulphates.quantile(0.98)
print(su_01)
print(su_98)
```

```
0.44
0.9
```

```
[75]: df=df[(df.sulphates>=su_01) & (df.sulphates<=su_98)]
sns.boxplot(df.sulphates)
```

```
[75]: <Axes: >
```




```
[76]: # Removing outliers from alcohol column
```

```
a_01=df.alcohol.quantile(0.01)
a_99=df.alcohol.quantile(0.99)
print(a_01)
print(a_99)
```

```
9.0
12.724
```

```
[77]: df=df[(df.alcohol>=a_01) & (df.alcohol<=a_99)]
sns.boxplot(df.alcohol)
```

```
[77]: <Axes: >
```



Therefore all the outliers are removed

0.0.3 Task - 3 : Machine Learning Model Building

```
[233]: # split into X and y
```

```
X =df.iloc[:, :-1]
X.head()
```

```
[233]: fixed_acidity volatile_acidity citric_acid residual_sugar chlorides \
```

```
0          7.4    0.70    0.00    1.9    0.076
```

```
1          7.8    0.88    0.00    2.6    0.098
```

```
2          7.8    0.76    0.04    2.3    0.092
```

```
3         11.2    0.28    0.56    1.9    0.075
```

```
4          7.4    0.70    0.00    1.9    0.076
```

```
free_sulfur_dioxide total_sulfur_dioxide density pH sulphates \
```

```
0          11.0          34.0 0.9978 3.51          0.56
```

```
1          25.0          67.0 0.9968 3.20          0.68
```

```
2          15.0          54.0 0.9970 3.26          0.65
```

3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4

```
[234]: Y =df.quality
        Y.head()
```

```
[234]: 0    5
```

```
1    5
```

```
2    5
```

```
3    6
```

```
4    5
```

```
Name: quality, dtype: int64
```

Label Binarisation (Conidering alcohol quality > 7 as good and assigning '1' to it else assigning '0')

```
[235]: Y = df['quality'].apply(lambda y_value: 1 if y_value>=7 else 0)
```

```
[236]: print(Y)
```

0	0
1	0
2	0
3	0
4	0
..	
1593	0
1594	0
1595	0
1596	0
1597	0

Name: quality, Length: 866, dtype: int64

[illegible]

```
[238]: X_train.shape
```

[238] : (692, 11)

```
[239]: X_test.shape
```

[239]: (174, 11)

```
[240]: print(Y.shape, Y_train.shape, Y_test.shape)
```

(866,) (692,) (174,)

0.0.4 Decision Tree Classifier

```
[242]: from sklearn.tree import DecisionTreeClassifier modell =  
DecisionTreeClassifier(max_depth=2, splitter='best', criterion='entropy'  
) modell.fit(X_train, Y_train)
```

```
[242]: DecisionTreeClassifier(criterion='entropy', max_depth=2)
```

```
[243]: d_y_predict = model1.predict(X_test)
        d_y_predict
```

[illegible]

```
[245]: d y predict train = model1.predict(X train)
```

0.0.5 Task - 4 : Evaluating the model (Decision tree classifier)

```
[246]: from sklearn.metrics import _
```



```

accuracy_score, classification_report, confusion_matrix
print('Testing Accuracy = ',
accuracy_score(Y_test,d_y_predict)) print('Training Accuracy =
', accuracy_score(Y_train,d_y_predict_train))

```

Testing Accuracy = 0.8793103448275862
Training Accuracy = 0.8916184971098265

0.0.6 Random Forest Classifier

```

[247]: from sklearn.ensemble import RandomForestClassifier model2
=RandomForestClassifier(n_estimators=200,criterion='entropy'
y') model2.fit(X_train,Y_train)

```

```

[247]: RandomForestClassifier(criterion='entropy', n_estimators=200)
[248]: r_y_predict = model2.predict(X_test)
r_y_predict_train = model2.predict(X_train)

```

0.0.7 Task - 4 : Evaluating Random Forest Model

```

[249]: print('Testing Accuracy = ', accuracy_score(Y_test,r_y_predict))
print('Training Accuracy = ',
accuracy_score(Y_train,r_y_predict_train))

```

Testing Accuracy = 0.9425287356321839
Training Accuracy = 1.0

0.0.8 Naive Bayesian Classification Model

```

[251]: from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train,Y_train)

```

```

[251]: GaussianNB()

```

```

[252]: y_pred2 = gnb.predict(X_test)
y_pred2

```

```

[252]: array([1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1,

```

```

0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0,
1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 0, 0])

```

0.0.9 Task - 4 : Evaluating Naive Bayesian Classification Model

```

[254]: from sklearn.metrics import accuracy_score
gnb_acc=accuracy_score(Y_test,y_pred2)
gnb_acc

```

```

[254]: 0.8850574712643678

```

0.1 Accuracies of all the algorithms used in model building phase :

Decision Tree Classification : 87.93 %

0.1.1 Random Forset Classification : 94.25 %

Naive Bayesian Classification : 88.50 %

0.1.2 Conclusion : Random Forest Classifier Model is best suited for the wine quality dataset.

0.1.3 Task - 5 : Test with random observation

```

[262]: input_data = [7.9, 1.0, 0, 3.0, 0.08, 30, 100, 0.9562, 3.1,
0.74, 11.5] prediction = model1.predict([input_data])
prediction

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
DecisionTreeClassifier was fitted with feature names
warnings.warn(

```

```

[262]: array([0])

```

According to “decision tree classifier” model, the above random observation gives prediction “array([0])” i.e., bad quality alcohol

```

[263]: input_data_2 = [7.9, 1.0, 0, 3.0, 0.08, 30, 100, 0.9562, 3.1,
0.74, 11.5] prediction2 = model2.predict([input_data_2])
prediction2

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
RandomForestClassifier was fitted with feature names
warnings.warn(

```

```
[263]: array([0])
```

According to “Random Forest classifier” model, the above random observation gives prediction “array([0])” i.e., bad quality alcohol

```
[264]: input_data_3 = [7.9, 1.0, 0, 3.0, 0.08, 30, 100, 0.9562, 3.1,
0.74, 11.5] prediction3 = gnb.predict([input_data_3])
prediction3
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but GaussianNB was
fitted with feature names
  warnings.warn(
```

```
[264]: array([0])
```

According to “Naive Bayesian classifier” model, the above random observation gives prediction “array([0])” i.e., bad quality alcohol

0.2 CONCLUSION : For the same random observation, all the three models gave the “alcohol quality is BAD”